



# CAR CRASH PREDICTION

"Modelo predictivo (Espacial) de  
sinistros en las calles de Santiago"

# "Modelo predictivo (Espacial) de siniestros en las calles de Santiago"

**UDD - Universidad del Desarrollo**

**MDS-18 BDA**

Peredo, Oscar

Professor

Gortari, Francisca

Rovai, Marcelo


Sacasa, Manuel

Master Candidates

# OBJECTIVE

---

The goal of the research is to predict a crash risk score for an urban grid with a 2013-2018 car crash georeferenced dataset and static urban descriptive public data set.



All the research was processed in MacBook Pro (Retina, 15-inch, 2017) with 2.9 GHz Intel Core i7 and 16GB 2133MHz LPDDR3 Memory, processing power on-premise machines.

The data science coding environment was **Anaconda's Jupyter notebooks** running Python 3.7.1 and PySpark 2.4.3. Maps, and Grids were developed mainly with GeoPandas 0.4.1, supported by several packages as Folium and OSMNX.

# PIPELINE

Main libraries used in the end to end pipeline Jupyter Notebook

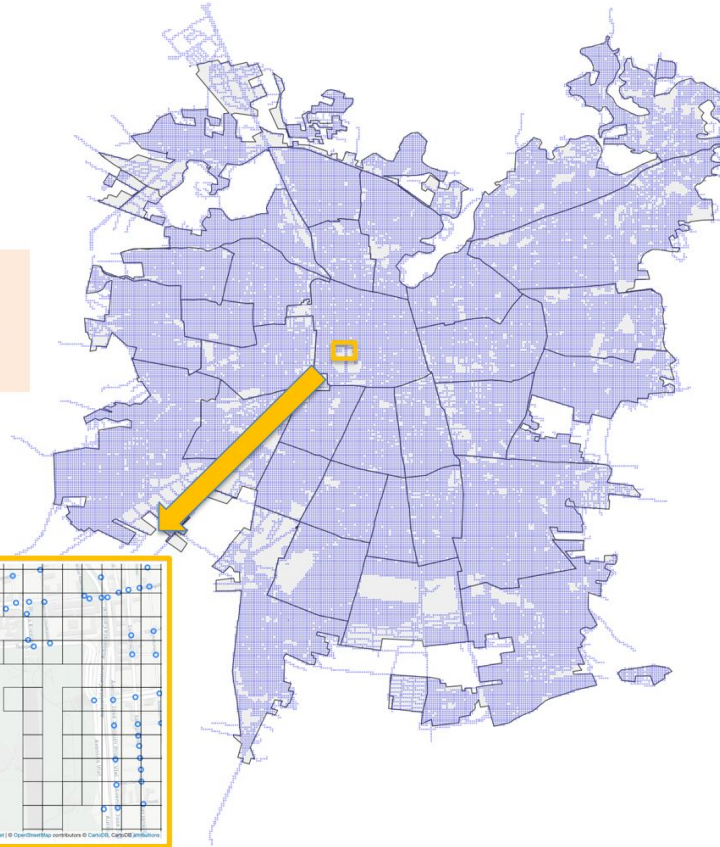

<b>OS</b>	<b>Geopandas</b>
<b>Math</b>	<b>Shapely</b>
<b>Numpy</b>	<b>MultiLineString</b>
<b>Pandas</b>	<b>Shapefile</b>
<b>Pylab</b>	<b>Gpd_lite_toolbox</b>
<b>Matplotlib</b>	<b>Findspark</b>
<b>Seaborn</b>	<b>Pyspark</b>
<b>Folium</b>	<b>Sparkxgb</b>
<b>Networkx</b>	
<b>Osmnx</b>	



# GRID

100 Meters Grid - Santiago Municipality

100m x 100m  
Grid to Road (CD)  
- 63,000 "cells"



A function was created to generate grids associated to roads. The input parameters are the length of each cell and the type of roads (or streets) to be considered. For this work we used 100m and all streets except the road fclasses "Service" and "Footway".

Also a 50m grid was generated for future work and it is available on grid folder.

# Filtering Roads

Filtering roads is important, in order to reduce the number of "cells" to be analysed by final model, once we will work only with cells where a car accident can really happen.




Roads with Service and Footway Filter

# DATA

## Static Features Dataset

The static features were obtained from [OPEN STREET MAP OSM CHILE](#) dataset with georeferenced points: **places, points of interest, traffic, transport, roads and intersections** (created from roads).



# DATA

## Static Features Dataset

### Creating a Function

```
1 def adding_static_feat_dist_to_grid(g,
2                                     osm_shp,
3                                     feature,
4                                     meters,
5                                     agg_type='sum'):
6     pois_x = osm_shp[osm_shp.fclass == feature]
7     x = pois_x[['geometry']].copy()
8     dist = (meters * 0.1) / 11000
9     x['geometry'] = x.geometry.buffer(dist)
10    x = gpd.sjoin(g, x, how='left', op='intersects')
11    x = x.fillna(0)
12    feature = feature + '_' + str(meters)
13    x = x.rename(columns={'index_right': feature})
14    x[feature] = x[feature].apply(lambda x: 0 if x == 0.0 else 1)
15    x = x.groupby('FID', as_index=False).agg({
16        feature: agg_type,
17        'geometry': 'first',
18    })
19    x = gpd.GeoDataFrame(x, crs='4326')
20    print("grid shape: ", x.shape)
21    print("Static Feature type {} has {} events".format(
22        feature, x[feature].sum()))
23    return x
```

executed in 6ms, finished 16:14:49 2019-08-02

```
1 g = grid
2 osm_shp = pois
3 meters = 100
```

executed in 3ms, finished 16:14:54 2019-08-02

```
1 feature = 'school'
2 school_100 = adding_static_feat_dist_to_grid(g, osm_shp, feature, meters)
```

executed in 6.29s, finished 16:15:06 2019-08-02

```
grid shape: (63029, 3)
Static Feature type school_100 has 12675 events
```

X	63029	non-null	float64
Y	63029	non-null	float64
bank	63029	non-null	int64
bench	63029	non-null	int64
beverages	63029	non-null	int64
bus_stop	63029	non-null	int64
bus_stop_100	63029	non-null	int64
cafe	63029	non-null	int64
convenience	63029	non-null	int64
convenience_100	63029	non-null	int64
convenience_200	63029	non-null	int64
crossing	63029	non-null	int64
crossing_100	63029	non-null	int64
fast_food	63029	non-null	int64
fast_food_100	63029	non-null	int64
fast_food_200	63029	non-null	int64
fuel	63029	non-null	int64
intercort	63029	non-null	int64
kindergarten	63029	non-null	int64
motorway_junction	63029	non-null	int64
parking	63029	non-null	int64
parking_bicycle	63029	non-null	int64
pharmacy	63029	non-null	int64
railway_station	63029	non-null	int64
railway_station_100	63029	non-null	int64
restaurant	63029	non-null	int64
restaurant_100	63029	non-null	int64
school	63029	non-null	int64
school_100	63029	non-null	int64
school_200	63029	non-null	int64
stop	63029	non-null	int64
stop_100	63029	non-null	int64
taxi	63029	non-null	int64
traffic_signals	63029	non-null	int64
traffic_signals_100	63029	non-null	int64
turning_circle	63029	non-null	int64

dtypes: float64(2), int64(34)



Grid individual Centroids







Static Features from OSM



# DATA

## CONASET

### 2013-18



# DATA

## Dynamic Features Dataset

The dynamic features were obtained from [CONASET](#)

6 Datasets: Georeferenced car crashes from 2013 until 2018 enriched with injury, type of crash, wounded persons, etc. Name of the datasets **"Siniestros RM20XX"**.

The "Siniestros" **datasets from 2013 to 2017 have been used to train/test the models**, while the **2018 dataset has been used to validate** the best model (Dynamic 2018 features generated from 2017 dataset - 'Year-1').



# DATA

## Dynamic Features Dataset

Car Crashes 2013-18 Santiago Municipality



2013 to 2018 Crashes by Sensus Zone - Santiago Municipality



# DATA

## Dynamic Features Dataset

Dataset concatenated  
Only Spatial data (no time)

```
1 siniestros.sample(10)
```

executed in 16ms, finished 17:35:27 2019-07-29

	geometry	Ano	Tipo_CONA	Fallecidos	Graves	Menos_Grav	Leves
18149	POINT (-70.57345845594295 -33.53564062681347)	2013	COLISION	0	0	0	1
56367	POINT (-70.73783930391549 -33.46285170610612)	2015	ATROPELLO	0	0	0	0
108348	POINT (-70.692421 -33.431526)	2018	COLISION	0	0	0	0
19879	POINT (-70.66134830647911 -33.55932925604031)	2013	COLISION	0	0	0	2
13145	POINT (-70.61845651655494 -33.54643059019953)	2013	COLISION	0	0	0	1
74735	POINT (-70.59810638248351 -33.4397427083309)	2016	COLISION	0	0	0	0
32657	POINT (-70.53339396042094 -33.47330588298253)	2014	ATROPELLO	0	0	0	1
90799	POINT (-70.7298924 -33.47838680000002)	2017	COLISION	0	0	0	1
6654	POINT (-70.6365284731105 -33.43756410420003)	2013	CAIDA	0	0	0	0
68232	POINT (-70.6113239 -33.4227138)	2016	CHOQUE	0	0	0	0

### 4.3 Adding Crash Severity

Severity of a Crash Definition: "The severity of a crash. Possible values are 'F' (Fallecidos), 'G' (Grave), 'M' (Menos Grave), 'L' (Leves), 'N' (non-injury). This is determined by the worst injury sustained in the crash at time of entry."

```
1 def sev_crash(row):
2     if row['Fallecidos'] != 0: return 'F'
3     elif row['Graves'] != 0: return 'G'
4     elif row['Menos_Grav'] != 0: return 'M'
5     elif row['Leves'] != 0: return 'L'
6     else: return 'N'
```

executed in 3ms, finished 17:35:51 2019-07-29

```
1 siniestros['SEV'] = siniestros.apply(sev_crash, axis=1)
```

executed in 3.09s, finished 17:36:06 2019-07-29

Creating a new Column Severity Index

```
1 def sev_index_crash(row):
2     if row['Fallecidos'] != 0: return 5
3     elif row['Graves'] != 0: return 4
4     elif row['Menos_Grav'] != 0: return 3
5     elif row['Leves'] != 0: return 2
6     else: return 1
```

executed in 4ms, finished 17:36:17 2019-07-29

```
1 siniestros['SEV_Index'] = siniestros.apply(sev_index_crash, axis=1)
```

executed in 3.32s, finished 17:36:32 2019-07-29


Adding Crash Severity



# DATA

## Dynamic Features with events from "meters"

Example event X meters from GRID [100m]



Function to capture event X meters from GRID

```
1 def adding_sin_type_date_dist_to_grid(g, siniestros, Y, D1, D2, meters):
2     s = siniestros[(siniestros.Fecha >= D1) & (siniestros.Fecha <= D2)]
3     s = s[s.Tipo_CONA == Y]
4     s = s[['geometry']].copy()
5     dist = (meters*0.1)/11000
6     s['geometry'] = s.geometry.buffer(dist)
7     s = gpd.sjoin(g, s, how='left', op='intersects')
8     s = s.fillna(0)
9     Y = Y + '_' + str(meters)
10    s = s.rename(columns={'index_right': Y})
11    s[Y] = s[Y].apply(lambda x: 0 if x == 0.0 else 1)
12    gs = s.groupby('FID', as_index=False).agg({
13        Y: 'sum',
14        'geometry': 'first'
15    })
16    gs = gpd.GeoDataFrame(gs, crs='4326')
17    print("Grid shape: ", gs.shape)
18    print("Crash type {} has {} events".format(Y, gs[Y].sum()))
19    return gs
```

executed in 8ms, finished 14:42:51 2019-07-26

```
1 D1 = '2017-01-01'
2 D2 = '2017-12-31'
3 Y = 'INCENDIO'
4 meters = 300
```

executed in 0.00s, finished 14:42:51 2019-07-26

```
1 = INCENDIO_300 = adding_sin_type_date_dist_to_grid(g, siniestros, Y, D1, D2,
2     meters)
```

executed in 0.17s, finished 14:42:51 2019-07-26

```
Grid shape: (87741, 3)
```

Crash type INCENDIO\_300 has 128 events

128

```
1 meters = 100
```

executed in 0.00s, finished 14:42:51 2019-07-26


```
1 = INCENDIO_100 = adding_sin_type_date_dist_to_grid(g, siniestros, Y, D1, D2,
2     meters)
```

executed in 0.52s, finished 14:42:51 2019-07-26

```
Grid shape: (87741, 3)
```

Crash type INCENDIO\_100 has 49 events

49



# DATA

## Dynamic Features and Dependent Variable

<del>ATROPELLO</del>	63029	non-null	int64
ATROPELLO_100	63029	non-null	int64
ATROPELLO_200	63029	non-null	int64
<del>CAIDA</del>	63029	non-null	int64
CAIDA_100	63029	non-null	int64
CAIDA_200	63029	non-null	int64
<del>CHOQUE</del>	63029	non-null	int64
CHOQUE_100	63029	non-null	int64
CHOQUE_200	63029	non-null	int64
<del>COLISION</del>	63029	non-null	int64
COLISION_100	63029	non-null	int64
COLISION_200	63029	non-null	int64
<del>FID</del>	63029	non-null	int64
<del>INCENDIO</del>	63029	non-null	int64
INCENDIO_100	63029	non-null	int64
INCENDIO_200	63029	non-null	int64
<del>OTRO TIPO</del>	63029	non-null	int64
OTRO TIPO_100	63029	non-null	int64
OTRO TIPO_200	63029	non-null	int64
<del>SEV_Index_1</del>	63029	non-null	float64
SEV_Index_100	63029	non-null	float64
SEV_Index_200	63029	non-null	float64
<del>VOLCADURA</del>	63029	non-null	int64
VOLCADURA_100	63029	non-null	int64
VOLCADURA_200	63029	non-null	int64

```
1 - def create_dep_var(row):
2 -     if (row['ATROPELLO'] + row['CAIDA'] + row['COLISION'] + row['INCENDIO'] +
3         row['OTRO TIPO'] + row['VOLCADURA']) == 0:
4         return 0
5     else:
6         return 1
```

executed in 4ms, finished 19:17:30 2019-07-29

```
1 train_dyna_features['SINIESTRO'] = train_dyna_features.apply(create_dep_var, axis=1)
2 test_dyna_features['SINIESTRO'] = test_dyna_features.apply(create_dep_var, axis=1)
```

executed in 5.41s, finished 19:18:45 2019-07-29

# DATA

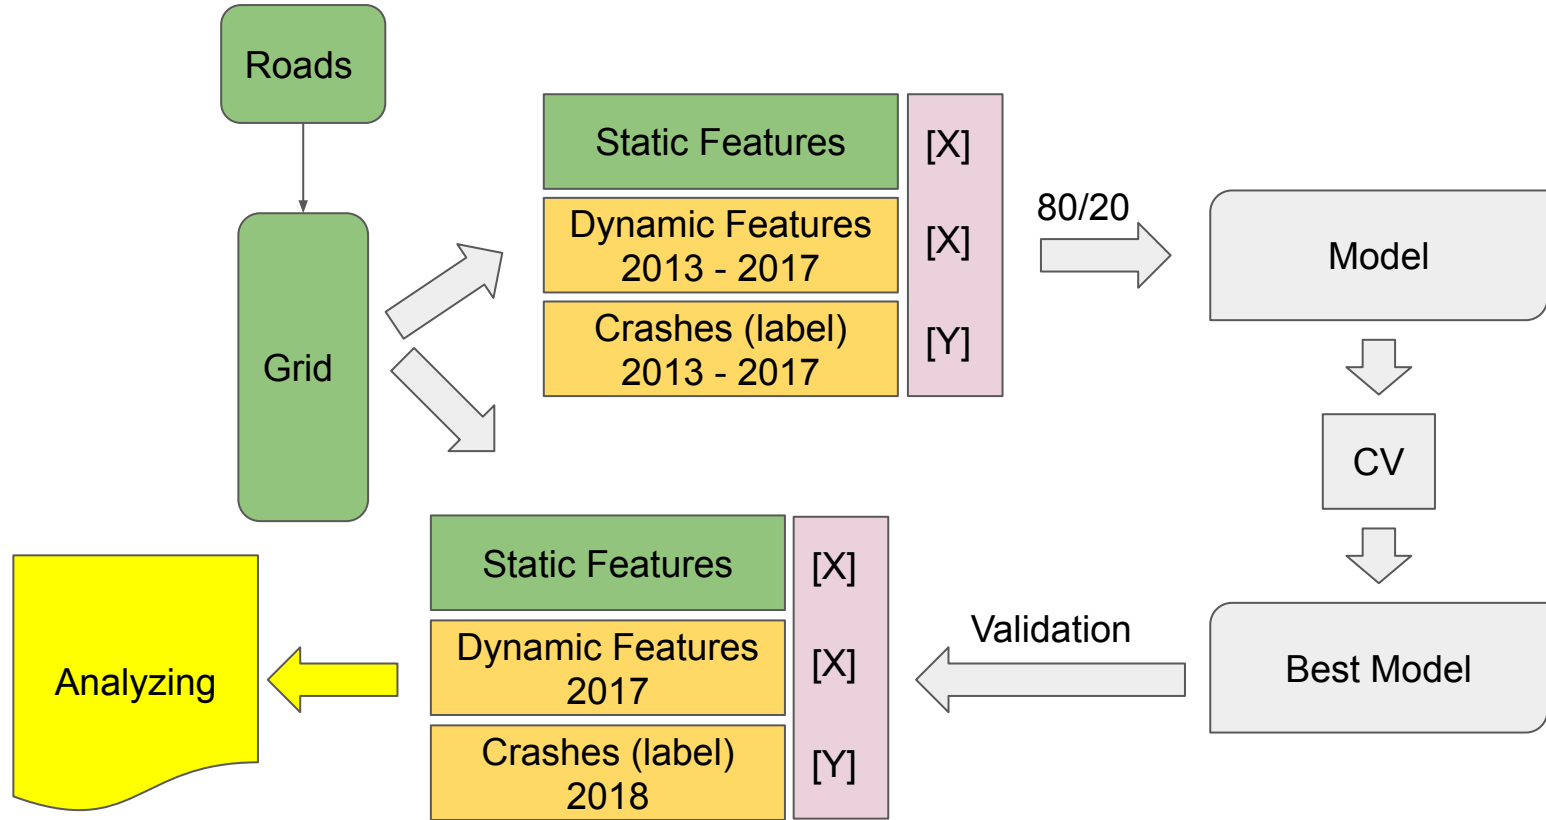
## Final datasets for modelling

A dataset called `final_train_dataset_grid_100.csv`, was used for training (80%) and test (20%) models. Here, dependent variable "SINIESTRO", was generated with all real events from 2013 to 2017. Same as its dynamic features.

A dataset called `final_test_dataset_grid_100.csv`, was used for validating the best model after cross validation phase. Here, dependent variable "SINIESTRO", was generated with all real events from 2018 but its dynamic features was from 2017 only.

```
root
|-- _c0: integer (nullable = true)
|-- id: integer (nullable = true)
|-- X: double (nullable = true)
|-- Y: double (nullable = true)
|-- bank: integer (nullable = true)
|-- bench: integer (nullable = true)
|-- beverages: integer (nullable = true)
|-- bus_stop: integer (nullable = true)
|-- bus_stop_100: integer (nullable = true)
|-- cafe: integer (nullable = true)
|-- convenience: integer (nullable = true)
|-- convenience_100: integer (nullable = true)
|-- convenience_200: integer (nullable = true)
|-- crossing: integer (nullable = true)
|-- crossing_100: integer (nullable = true)
|-- fast_food: integer (nullable = true)
|-- fast_food_100: integer (nullable = true)
|-- fast_food_200: integer (nullable = true)
|-- fuel: integer (nullable = true)
|-- intercert: integer (nullable = true)
|-- kindergarten: integer (nullable = true)
|-- motorway_junction: integer (nullable = true)
|-- parking: integer (nullable = true)
|-- parking_bicycle: integer (nullable = true)
|-- pharmacy: integer (nullable = true)
|-- railway_station: integer (nullable = true)
|-- railway_station_100: integer (nullable = true)
|-- restaurant: integer (nullable = true)
|-- restaurant_100: integer (nullable = true)
|-- school: integer (nullable = true)
|-- school_100: integer (nullable = true)
|-- school_200: integer (nullable = true)
|-- stop: integer (nullable = true)
|-- stop_100: integer (nullable = true)
|-- taxi: integer (nullable = true)
|-- traffic_signals: integer (nullable = true)
|-- traffic_signals_100: integer (nullable = true)
|-- turning_circle: integer (nullable = true)
|-- ATROPELLO_100: integer (nullable = true)
|-- ATROPELLO_200: integer (nullable = true)
|-- CAIDA_100: integer (nullable = true)
|-- CAIDA_200: integer (nullable = true)
|-- CHOQUE_100: integer (nullable = true)
|-- CHOQUE_200: integer (nullable = true)
|-- COLISION_100: integer (nullable = true)
|-- COLISION_200: integer (nullable = true)
|-- INCENDIO_100: integer (nullable = true)
|-- INCENDIO_200: integer (nullable = true)
|-- OTRO TIPO_100: integer (nullable = true)
|-- OTRO TIPO_200: integer (nullable = true)
|-- SEV_Index_100: double (nullable = true)
|-- SEV_Index_200: double (nullable = true)
|-- VOLCADURA_100: integer (nullable = true)
|-- VOLCADURA_200: integer (nullable = true)
|-- SINIESTRO: integer (nullable = true)
```

# TASKS





# MODELING

Pipeline, Split data and  
RF modelling as a baseline

```
1 categoricalColumns = cat_cols
2 cols = df.columns
3 stages = []
4
5 for categoricalCol in categoricalColumns:
6     stringIndexer = StringIndexer(inputCol=categoricalCol,
7                                     outputCol=categoricalCol + 'Index')
8     encoder = OneHotEncoderEstimator(inputCols=[stringIndexer.getOutputCol()],
9                                       outputCols=[categoricalCol + "classVec"])
10    stages += [stringIndexer, encoder]
11
12 label_stringIdx = StringIndexer(inputCol='SINIESTRO', outputCol='label')
13 stages += [label_stringIdx]
14
15 # Assemble the columns into a feature vector
16 assemblerInputs = [c + "classVec" for c in categoricalColumns] + numericCols
17 assembler = VectorAssembler(inputCols=assemblerInputs, outputCol="features")
18 stages += [assembler]
```

executed in 13ms, finished 12:14:21 2019-08-03

```
1 pipeline = Pipeline(stages = stages)
2 pipelineModel = pipeline.fit(df)
3 df = pipelineModel.transform(df)
4 selectedCols = ['label', 'features'] + cols
5 df = df.select(selectedCols)
6 df.printSchema()
```

executed in 291ms, finished 12:14:22 2019-08-03

```
1 train, test = df.randomSplit([0.8, 0.2], seed=24)
```

executed in 12ms, finished 12:14:23 2019-08-03

```
1 rf = RandomForestClassifier(featuresCol='features', labelCol='label')
2 rfModel = rf.fit(train)
3 predictions = rfModel.transform(test)
```

executed in 2.70s, finished 12:14:27 2019-08-03

```
1 predictions.select('id', 'label', 'rawPrediction', 'prediction',
2                    'probability').show(10)
```

executed in 477ms, finished 12:14:30 2019-08-03

id	label	rawPrediction	prediction	probability
0	0.0	[19.3985287329805...	0.0	[0.96992643664902...
3	0.0	[19.3985287329805...	0.0	[0.96992643664902...
14	0.0	[19.3985287329805...	0.0	[0.96992643664902...
18	0.0	[19.3985287329805...	0.0	[0.96992643664902...
23	0.0	[19.3985287329805...	0.0	[0.96992643664902...
33	0.0	[19.3985287329805...	0.0	[0.96992643664902...
34	0.0	[19.3985287329805...	0.0	[0.96992643664902...
44	0.0	[19.3985287329805...	0.0	[0.96992643664902...
56	0.0	[19.3985287329805...	0.0	[0.96992643664902...
54	0.0	[19.3985287329805...	0.0	[0.96992643664902...

only showing top 10 rows

```
1 evaluator = BinaryClassificationEvaluator()
2 print("Test Area Under ROC: " + str(
3     evaluator.evaluate(predictions, {evaluator.metricName: "areaUnderROC"})))
```

executed in 617ms, finished 12:14:43 2019-08-03

Test Area Under ROC: 0.8294667464568877

```
1 evaluator = MulticlassClassificationEvaluator()
2 accuracy = evaluator.evaluate(predictions, {evaluator.metricName: "accuracy"})
3 print("Accuracy: " + str(accuracy))
```

executed in 798ms, finished 12:14:57 2019-08-03

Accuracy: 0.7782577959311916

```
1 print("Test Error = %g" % (1.0 - accuracy))
```

executed in 3ms, finished 12:15:01 2019-08-03

Test Error = 0.221742

# MODELING

## GBT modelling and Feature Importance

idx	name	score
42	COLISION_100	0.439067
11	crossing	0.090049
17	intercect	0.071686
5	bus_stop	0.055971
33	traffic_signals	0.047034
36	ATROPELLO_100	0.035349
1	Y	0.028571
6	bus_stop_100	0.023284
34	traffic_signals_100	0.021061
48	SEV_Index_100	0.016812
26	restaurant_100	0.016779
12	crossing_100	0.015766
25	restaurant	0.015591
46	OTRO TIPO_100	0.014711
50	VOLCADURA_100	0.013828
22	pharmacy	0.013157
43	COLISION_200	0.012653
0	X	0.011431
41	CHOQUE_200	0.009117
38	CAIDA_100	0.008853

```
1 gbt = GBTCClassifier(maxIter=10)
```

executed in 7ms, finished 12:15:07 2019-08-03

```
1 gbtModel = gbt.fit(train)
```

executed in 7.27s, finished 12:15:15 2019-08-03

```
1 predictions = gbtModel.transform(test)
```

executed in 48ms, finished 12:15:17 2019-08-03

```
1 predictions.select('id', 'label', 'rawPrediction', 'prediction',  
2 'probability').show(5)
```

executed in 454ms, finished 12:15:19 2019-08-03

	id	label	rawPrediction	prediction	probability
0	0.0	[1.32412711336971...	0.0	[0.93390330915520...	
3	0.0	[1.32412711336971...	0.0	[0.93390330915520...	
14	0.0	[1.32412711336971...	0.0	[0.93390330915520...	
18	0.0	[1.32412711336971...	0.0	[0.93390330915520...	
23	0.0	[1.32412711336971...	0.0	[0.93390330915520...	

only showing top 5 rows

```
1 evaluator = BinaryClassificationEvaluator()  
2 print("Test Area Under ROC: " + str(  
3 evaluator.evaluate(predictions,  
4 {evaluator.metricName: "areaUnderROC"})))
```

executed in 542ms, finished 12:15:22 2019-08-03

Test Area Under ROC: 0.8383690618564904

```
1 evaluator = MulticlassClassificationEvaluator()  
2 accuracy = evaluator.evaluate(predictions, {evaluator.metricName: "accuracy"})  
3 print("Accuracy: " + str(accuracy))
```

executed in 720ms, finished 12:15:29 2019-08-03

Accuracy: 0.7879978006440971

```
1 print("Test Error = %g" % (1.0 - accuracy))
```

executed in 4ms, finished 12:15:30 2019-08-03

Test Error = 0.212002

```

1 ▾ xgboost = XGBoostEstimator(
2     featuresCol="features",
3     labelCol="label",
4     predictionCol="prediction"
5 )

```

executed in 58ms, finished 13:35:42 2019-08-03

```

1 xgbModel = xgboost.fit(train)

```

executed in 2.83s, finished 13:35:48 2019-08-03

```

1 predictions = xgbModel.transform(test)

```

executed in 298ms, finished 13:35:49 2019-08-03

```

1 predictions.select('id', 'label', 'prediction').show(5)

```

executed in 806ms, finished 13:35:50 2019-08-03

```

+---+-----+-----+
| id|label|prediction|
+---+-----+-----+
| 0| 0.0|      0.0|
| 3| 0.0|      0.0|
|14| 0.0|      0.0|
|18| 0.0|      0.0|
|23| 0.0|      0.0|
+---+-----+-----+
only showing top 5 rows

```

```

1 evaluator = MulticlassClassificationEvaluator(labelCol='label', metricName="accuracy")

```

executed in 9ms, finished 13:35:54 2019-08-03

```

1 accuracy = evaluator.evaluate(predictions)
2 print("Test Error = %g" % (1.0 - accuracy))

```

executed in 1.79s, finished 13:35:57 2019-08-03

Test Error = 0.210117

Almost same as Test Error (0.212002) got with GBT Default parameters.

# MODELING

## XGBoost

# MODELING

## Tuning and saving best GBT Model

```
1 predictions.select('id', 'label', 'rawPrediction', 'probability', 'prediction').show(10)
executed in 502ms, finished 12:26:23 2019-08-03
```

id	label	rawPrediction	probability	prediction
0	0.0	[1.54341621871724...	[0.95634631650520...	0.0
3	0.0	[1.54341621871724...	[0.95634631650520...	0.0
14	0.0	[1.54341621871724...	[0.95634631650520...	0.0
18	0.0	[1.54341621871724...	[0.95634631650520...	0.0
23	0.0	[1.54341621871724...	[0.95634631650520...	0.0
33	0.0	[1.54341621871724...	[0.95634631650520...	0.0
34	0.0	[1.54341621871724...	[0.95634631650520...	0.0
44	0.0	[1.54341621871724...	[0.95634631650520...	0.0
56	0.0	[1.54341621871724...	[0.95634631650520...	0.0
54	0.0	[1.54341621871724...	[0.95634631650520...	0.0

only showing top 10 rows

```
1 bestModel = cvModel.bestModel
executed in 3ms, finished 12:26:25 2019-08-03
```

```
1 bestModel
executed in 5ms, finished 12:26:26 2019-08-03
```

GBTClassificationModel (uid=GBTClassifier\_95ac8c178565) with 20 trees

```
1 bestModel.write().overwrite().save('./../model/GeoProjectBestModel_1.model')
executed in 721ms, finished 12:26:28 2019-08-03
```

```
1 paramGrid = (ParamGridBuilder()
2             .addGrid(gbt.maxDepth, [4, 6])
3             .addGrid(gbt.maxBins, [40, 60, 70])
4             .addGrid(gbt.maxIter, [10, 20])
5             .build())
6
7 cv = CrossValidator(estimator=gbt,
8                    estimatorParamMaps=paramGrid,
9                    evaluator=evaluator,
10                   numFolds=5)
11
12 # Run cross validations. This can take about 7.3 minutes!
13 cvModel = cv.fit(train)
14 predictions = cvModel.transform(test)
```

executed in 7m 15s, finished 12:23:56 2019-08-03

```
1 evaluator = BinaryClassificationEvaluator()
2 print("Test Area Under ROC: " + str(
3     evaluator.evaluate(predictions,
4                       {evaluator.metricName: "areaUnderROC"})))
```

executed in 537ms, finished 12:25:40 2019-08-03

Test Area Under ROC: 0.841362042678081

```
1 evaluator = MulticlassClassificationEvaluator()
2 accuracy = evaluator.evaluate(predictions, {evaluator.metricName: "accuracy"})
3 print("Accuracy: " + str(accuracy))
```

executed in 766ms, finished 12:25:42 2019-08-03

Accuracy: 0.7924750608750295

```
1 print("Test Error = %g" % (1.0 - accuracy))
```

executed in 3ms, finished 12:25:45 2019-08-03

Test Error = 0.207525



# VALIDATING

## Best GBT Model with 2018 dataset

```
1 valModel = GBTClassificationModel.load("../model/GeoProjectBestModel_1.model")
executed in 455ms, finished 12:39:11 2019-08-03
```

Loading the dataset with 2018 data:

Load Best Model

```
1 df_2018 = spark.read.csv('../data/final_test_dataset_grid_100.csv',
2                           header=True,
3                           inferSchema=True)
4 df_2018.printSchema()
```

executed in 308ms, finished 12:39:13 2019-08-03

Load Data

```
root
|-- _c0: integer (nullable = true)
|-- id: integer (nullable = true)
|-- X: double (nullable = true)
|-- Y: double (nullable = true)
```

```
1 num_cols, cat_cols = find_num_cat_features(df_2018)
```

executed in 4ms, finished 12:39:13 2019-08-03

0 categorical features  
54 numerical features

```
1 df_2018.groupby('SINIESTRO').count().toPandas()
```

executed in 216ms, finished 12:39:14 2019-08-03

SINIESTRO	count
0	1 6687
1	0 56342

Note that on this dataset, we have around 10.5% of the Grid's cells with crash events and 89.4% with No events.

```
1 numericCols = num_cols[1:-1] # Taking out id and Target variable
2 numericCols
```

executed in 5ms, finished 12:39:18 2019-08-03

```
['X',
 'Y',
 'bank',
 'bench',
 ..]
```

```
1 categoricalColumns = cat_cols
2 cols = df.columns
3 stages = []
4
5 for categoricalCol in categoricalColumns:
6     stringIndexer = StringIndexer(inputCol=categoricalCol,
7                                   outputCol=categoricalCol + 'Index')
8     encoder = OneHotEncoderEstimator(inputCols=[stringIndexer.getOutputCol()],
9                                     outputCols=[categoricalCol + "classVec"])
10    stages += [stringIndexer, encoder]
11
12 label_stringIdx = StringIndexer(inputCol='SINIESTRO', outputCol='label')
13 stages += [label_stringIdx]
14
15 assemblerInputs = [c + "classVec" for c in categoricalColumns] + numericCols
16 assembler = VectorAssembler(inputCols=assemblerInputs, outputCol='features')
17 stages += [assembler]
18
19 pipeline = Pipeline(stages = stages)
20 pipelineModel = pipeline.fit(df_2018)
21 df_2018 = pipelineModel.transform(df_2018)
22 selectedCols = ['label', 'features'] + cols
23 df_2018 = df_2018.select(selectedCols)
24 df_2018.printSchema()
```

executed in 242ms, finished 12:39:19 2019-08-03

Prepare Data

```
1 evaluator = BinaryClassificationEvaluator()
2 print("Test Area Under ROC: " + str(
3     evaluator.evaluate(predict_2018,
4                       {evaluator.metricName: "areaUnderROC"})))
```

executed in 994ms, finished 12:42:35 2019-08-03

Test Area Under ROC: 0.7742939707280319

```
1 evaluator = MulticlassClassificationEvaluator()
2 accuracy = evaluator.evaluate(predict_2018, {evaluator.metricName: "accuracy"})
3 print("Accuracy: " + str(accuracy))
```

executed in 981ms, finished 12:42:48 2019-08-03

Accuracy: 0.8942708911770773

```
1 print("Test Error = %g" % (1.0 - accuracy))
```


executed in 3ms, finished 12:40:51 2019-08-03

Test Error = 0.105729

# ANALYZING

## Results from Best Model

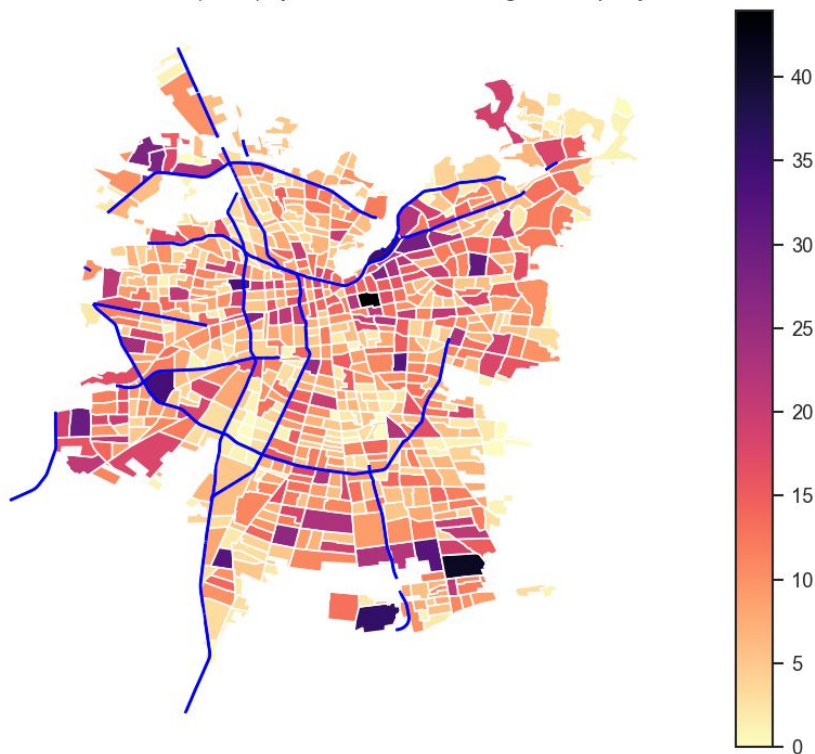
1	<code>print(classification_report(y_test,y_pred))</code>				
executed in 133ms, finished 13:14:48 2019-08-03					
	precision	recall	f1-score	support	
0.0	0.92	0.96	0.94	56342	
1.0	0.50	0.33	0.40	6687	
accuracy			0.89	63029	
macro avg	0.71	0.65	0.67	63029	
weighted avg	0.88	0.89	0.88	63029	



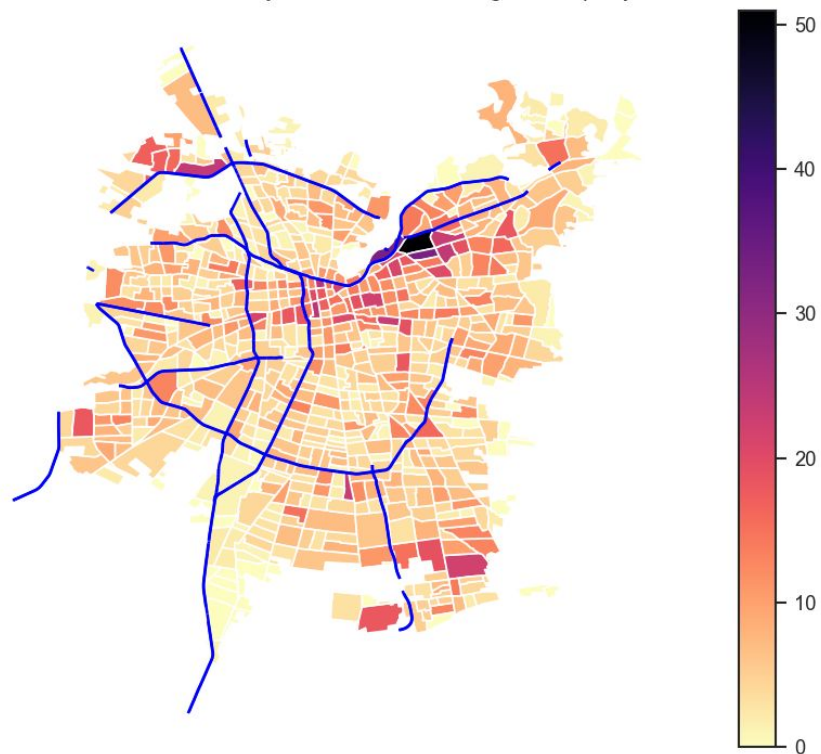
# ANALYZING

## Label versus Prediction

2018 Real Crashes (Label) by Sensus Zone - Santiago Municipality



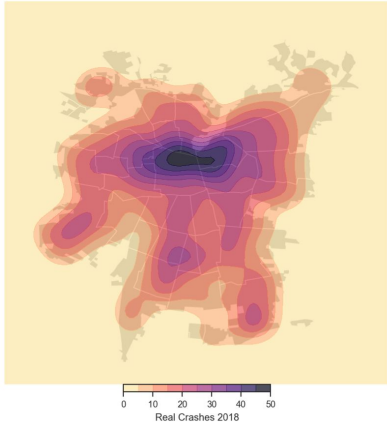
2018 Crash Predictions by Sensus Zone - Santiago Municipality







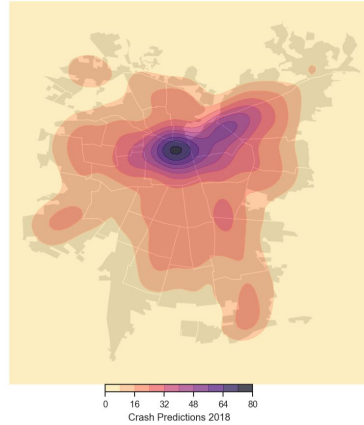
# ANALYZING

## Label versus Prediction

2018 Real Crashes (Label)



2018 Predicted Crashes





Proximity with  
highways/main roads  
seems to be an  
important factor of  
events (not well  
captured by model)



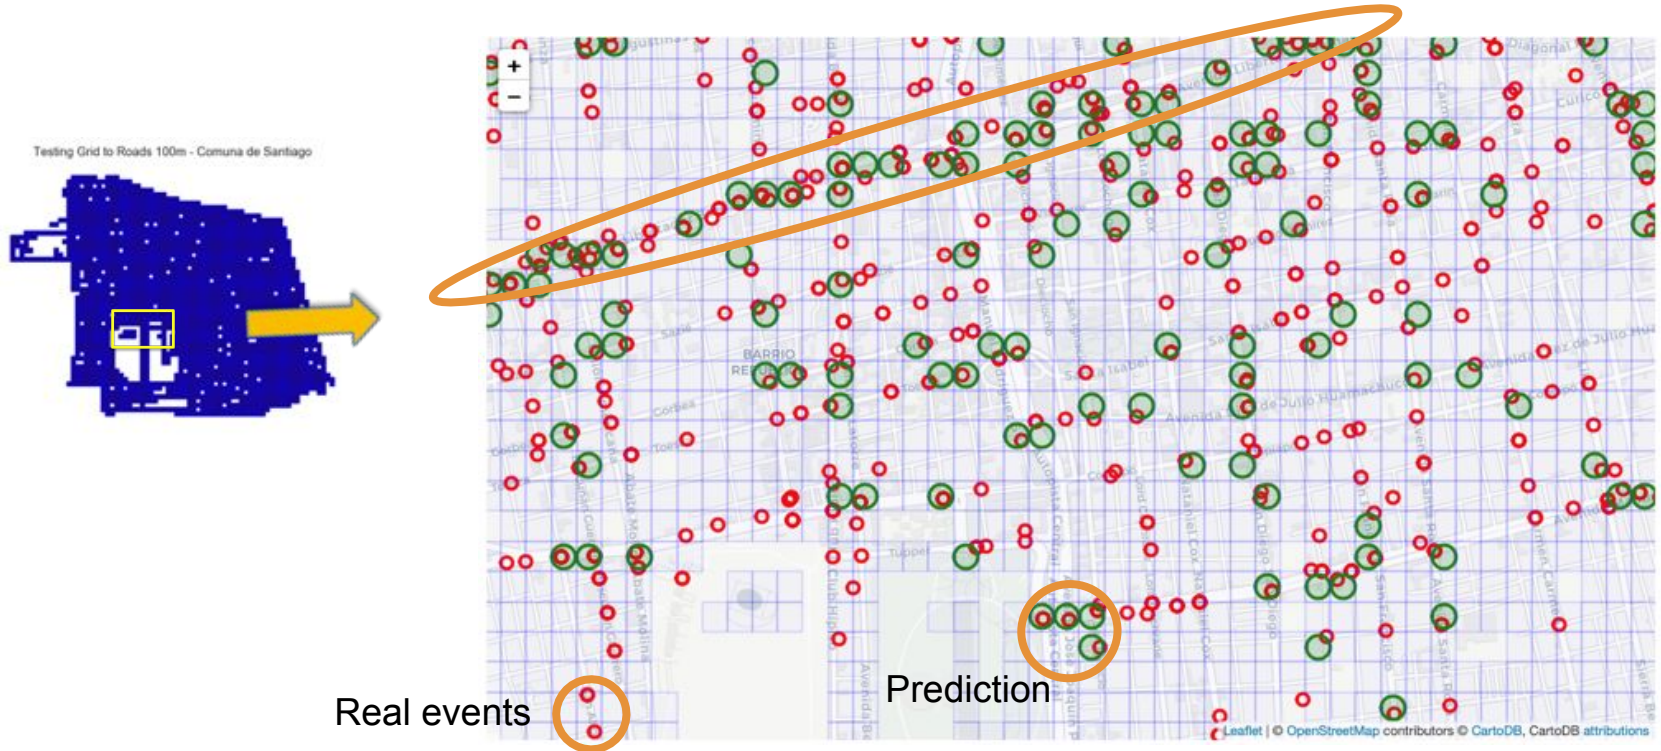
# ANALYZING

Label versus Main Roads

# ANALYZING

## Label versus Prediction

Alameda O'Higgins

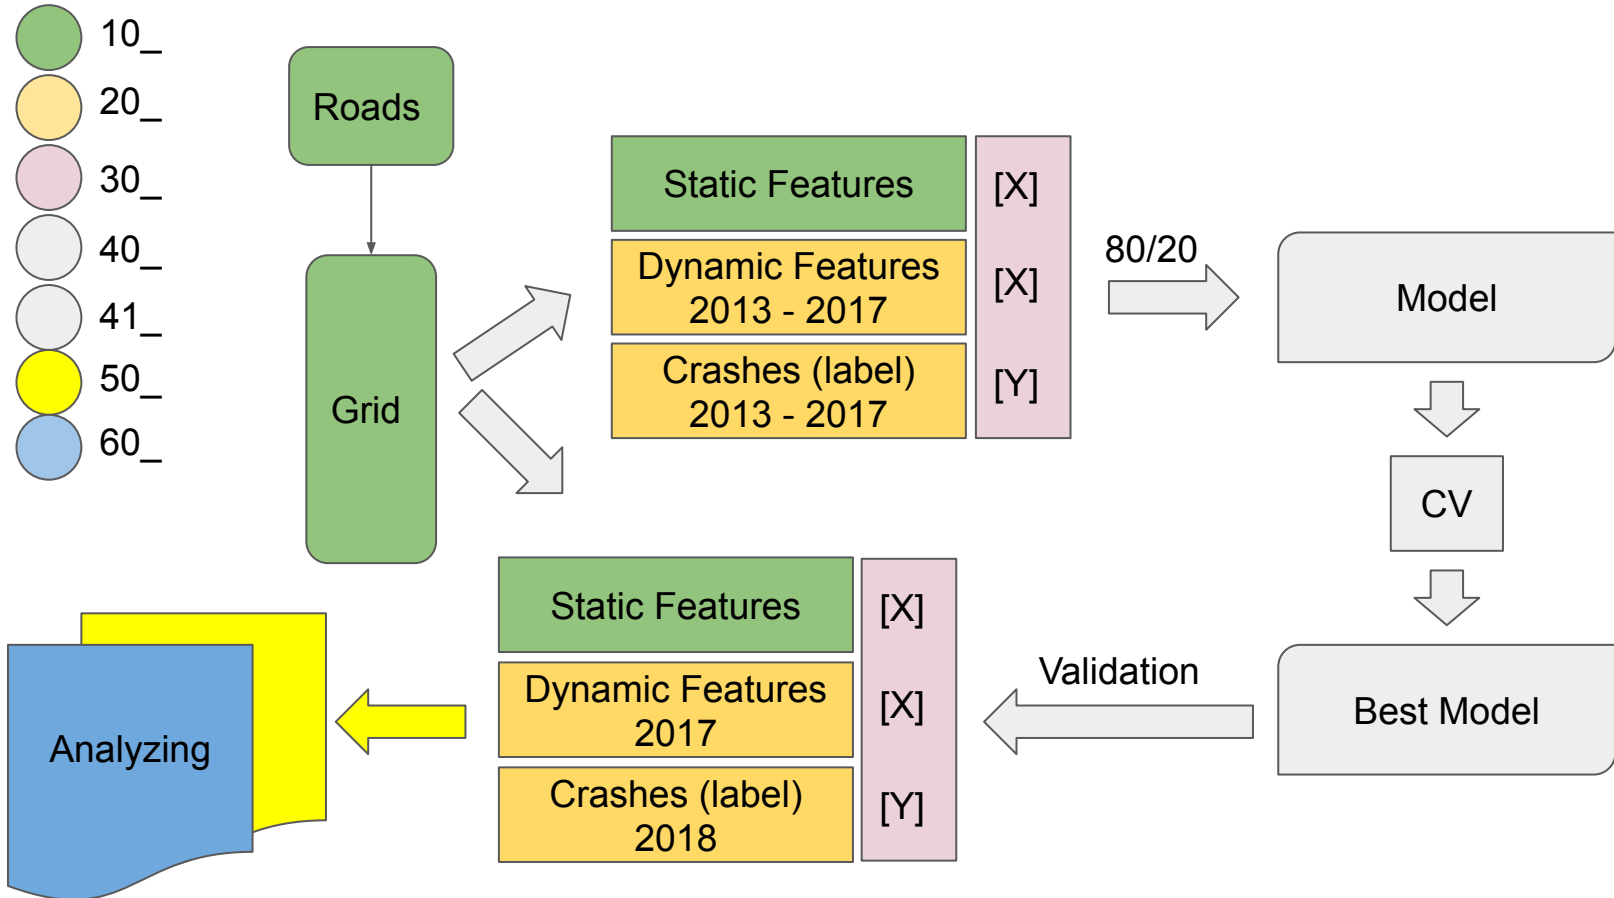


# FUTURE WORK

- Test model with different grids (50m and 200m for example)
- Form a dataset only with years that have more data, adding temporal dynamic features as working days, street conditions, clima, etc.
- Introduce new static features as:
  - Average road speed
  - Demographic density
  - Mobility factor
  - Roads geometry (curvature, inclination, etc.)
  - Altitude
  - Number of road lanes
  - Proximity with highways/main roads
  - Other

# Addendum

# NOTEBOOKS KEY





1	Datasets (Landing)	
1.1	CONASET	10_
1.2	OpenStreetMap	
2	Main Libraries	
3	Main Functions	
4	Getting Raw Data	
4.1	Restrict Analysis inside Santiago Municipality Area	
4.2	Getting POIs	
4.3	Getting Roads	
4.4	Getting Traffic Points	
4.5	Getting Transportation Points	
4.6	Getting Roads Intersections	
4.7	Visualizing a sample of Intersections (Santiago County)	
5	Creating a grid	
5.1	Grid related with Intersections	
5.2	Creating a generic grid creation function	
5.3	Testing generic Grid	
5.4	Grid related to Roads (Calles Discretizadas 'CD')	
5.5	Creating a 50 meters Grid to Roads	
5.6	Testing Grid 100 to roads	
5.7	Testing Grid 50 to roads	
6	Adding Static features to road grid ('CD')	
6.1	Including Intersections on Grid	
6.2	Join POI & Grid	
6.3	Join Traffic Points to Grid	
6.4	Join Transportation Points to Grid	
7	Adding new Static Features on a distance (meters)	
7.1	Creating a Function	
8	Creating a Static Features dataset	
8.1	Creating a .CSV Static Feature dataset	

# NOTEBOOKS

# CONTENT

1	Datasets (Landing)	
1.1	CONASET	
1.2	OpenStreetMap	20_
2	Main Libraries	
3	Main Functions	
4	Getting Raw Data	
4.1	Loading data from CONASET	
4.2	Restrict Siniestros inside Santiago Municipality Area	
4.3	Adding Crash Severity	
5	Filter crash type per year	
6	Adding Crash related features to grid	
6.1	Defining a generic Function to add a crash event by type and date range to the grid	
7	Define Features from 2013 to 2017 to be used as Train	
7.1	Defining Crash Type Buffer	
7.2	Defining a generic Function to add a crash event by type, distance in meters and date range to the grid	
7.3	Defining a generic Function to add a crash dinamic features and date range to the grid	
8	Creating a Train dataset	
8.1	Creating a .CSV Static Feature dataset	
9	Defining Features from 2018 to be used as a Test dataset	
9.1	Crash Events by type	
9.2	Crash Events per type and distance	
10	Creating a Test dataset	
10.1	Creating a .CSV Static Feature dataset	

# NOTEBOOKS CONTENT

30\_

- 1 Datasets (Landing)
  - 1.1 CONASET
  - 1.2 OpenStreetMap
- 2 Main Libraries
- 3 Creating Final Train and Test datasets
- 4 Creating and Saving Train Dataset
- 5 Creating and Saving Test Dataset

50\_

- 1 Datasets (Landing)
  - 1.1 CONASET
  - 1.2 OpenStreetMap
- 2 Main Libraries
- 3 Analyzing the Crash Prevision 2018 Dataset

40\_

- 1 Preliminar Installation
- 2 General functions
- 3 Libraries and pySpark initialization
- 4 Dataset - Generated from Notebook:
- 5 Verifying dataset balance
- 6 Preparing Data for Machine Learning
- 7 RF Model
- 8 GBT Model
- 9 Tuning The GBT Model
- 10 Validating Best Model with 2018 dataset

60\_

- 1 Datasets (Landing)
  - 1.1 CONASET
  - 1.2 OpenStreetMap
- 2 Main Libraries and Initialization
- 3 Main Functions
- 4 Import and visualize dataset
- 5 Model Result Spatial Visualization
- 6 Test Results on Santiago County
- 7 Working with "Sensal Zones"
  - 7.1 Create a map with Crashes and zones

41\_

- 1 Preliminar Installations
- 2 General functions
- 3 Libraries and pySpark initialization
- 4 Dataset - Generated from Notebook:
- 5 Verifying dataset balance
- 6 Preparing Data for Machine Learning
- 7 XGBoost Model



# CAR CRASH PREDICTION

"Modelo predictivo (Espacial) de  
siniestros en las calles de Santiago"